

I'm working on a wildfire spread model using geotiffs. The wind direction band is in degrees from north. Are there other ways to represent this? Like cardinal directions? Otherwise I think it will be difficult for the U-Net model to read.

ChatGPT said:

Yes — you should **not** feed raw wind direction in degrees (0–360) directly into a U-Net. Angular variables are *circular*, and treating them as linear scalars creates discontinuities (e.g., $359^\circ \approx 0^\circ$, but numerically they look far apart). This absolutely harms training.

Use one of the following standard encodings.

1. Sine–Cosine Encoding (recommended for ML models)

Convert direction θ (degrees clockwise from north) into two channels:

$$u = \sin(\theta \cdot \pi/180) \quad v = \cos(\theta \cdot \pi/180)$$

These form a unit vector with continuous topology around 360° .

Add both **u** and **v** as two separate raster bands.

Advantages

- Removes angular discontinuity
- Allows the U-Net to learn directional influence smoothly
- Keeps magnitude independent (so wind speed can be its own band)

This is the **standard encoding** in atmospheric science, oceanography, and robotics.

2. Wind Vector Encoding (direction + speed fused)

If you have wind speed S :

$$u = S \cdot \sin(\theta) \quad v = S \cdot \cos(\theta)$$

This is the meteorology convention: (u = eastward, v = northward wind components). Gives the model the actual physical vector.

3. Cardinal Direction Binning (not ideal for U-Net)

You can bin the direction into classes:

- N, NE, E, SE, S, SW, W, NW → 8 classes
or
- 16-point compass

Represent as a one-hot or ordinal raster.

Problems

- Loses detail
- Creates sharp artificial boundaries
- U-Net will treat it like a semantic segmentation label rather than physical variable

Only useful if your dataset is extremely noisy.

4. Complex Number Encoding (equivalent to sin/cos)

Represent direction as:

$$z = \cos(\theta) + i\sin(\theta) \quad z = \cos(\theta) + i\sin(\theta)$$

In practice this is identical to providing sin/cos as two bands.

5. Principal Component Embedding (overkill)

You can embed direction using a learned 2D circular embedding, but this provides no real advantage over sin/cos for physical data.

Recommendation for your wildfire model

Use **sin/cos encoding** or **u/v wind components**, depending on whether wind speed is already a band.

If you already have wind speed:

→ Use u/v wind components.

If you only know direction:

→ Use sin/cos as two channels.

This will materially improve stability and learning of directional influence on spread.

If you'd like, I can show you code to convert your GeoTIFF wind direction band into two new raster bands with sin/cos encoding.

You said:

this is a standard practice? site sources

ChatGPT said:

Yes — encoding wind direction using vector components (or sine / cosine) rather than raw degrees is standard practice. Here are a few references and examples from meteorology / wind-field modelling that justify common alternatives:



Established practice & scientific precedent

- In data assimilation systems used for weather forecasting (e.g. WRFDA) observed wind “speed + direction” is routinely converted into zonal (u) and meridional (v) components prior to assimilation. [Tellus Journals+1](#)
- In atmospheric / climate-analysis studies, authors note that because wind is fundamentally a **vector** quantity (not a scalar), representation by (u, v) components is more physically meaningful than discrete direction codes. [ASCMO+1](#)
- In machine-learning (CNN)-based wind / flow modelling, some recent works explicitly compare performance between a model using (u, v) or normalized sine/cosine of direction versus one using raw directional angles — and find that the sine/cosine (or normalized) encoding outperforms raw-angle encoding, especially for stability and

continuity of flow fields. [ScienceDirect+1](#)

Why vector-/sine-cosine encoding is the right approach

- Angular variables (0° – 360°) are *circular*. A small change near the 0° / 360° boundary (e.g. $359^\circ \rightarrow 1^\circ$) is actually a small shift in direction — but numerically that's a large jump if encoded as a scalar. That discontinuity tends to confuse models that treat variables linearly.
 - Converting direction θ and speed S into $u = S \cdot \cos(\theta)$, $v = S \cdot \sin(\theta)$ (or normalized sin/cos if speed is separate) preserves the continuity in direction and frames the wind as a continuous vector field. That is the same mathematical representation used in meteorology, oceanography, climate models, and many fluid-dynamics studies.
[Atmospheric and Oceanic Sciences+2UPCommons+2](#)
-

* Example from recent ML-wind-modelling literature

A recent paper on downscaling of surface-wind forecasts via convolutional neural networks describes two variants:

- One forecasting (u, v) components directly,
- Another forecasting normalized direction components (i.e. sine and cosine of direction, independent from speed),

and show that the normalized sine/cosine variant performs better — particularly in producing sharp, physically plausible flow features — than the raw (u, v) variant when speed errors affect direction accuracy. [NPG+1](#)

Conclusion for your wildfire-spread use case

Using sine/cosine encoding or wind vector components is not merely a heuristic — it is a well-accepted, physically grounded standard in meteorology, atmospheric modelling, climate

research, and ML-based flow prediction. Given the circular nature of wind direction, this encoding substantially reduces artifacts and helps models (like a U-Net) learn a physically consistent directional field.

If you like, I can pull up **5 to 10 peer-reviewed** references (with DOI) that specifically deal with vector encoding of wind (or other circular variables) in ML or physical-modelling contexts — which could be useful to cite in a publication.

ChatGPT c

You said:

im working with a variety of GIS files to create a dataset to feed into a unet. The unit takes tifs. what libraries and standard processes can i use to convert .shp files, geotiffs, and csv files to a tiff? can i store time series data in a tiff? like a time stamp with each band? How many bands can a geotiff have? provide sources for possible work flows and provided information.

ChatGPT said:

Here's a technical breakdown of how you can convert your various GIS data types (e.g. shapefiles, CSV, GeoTIFF) into TIFF/GeoTIFF input suitable for a UNet — and answers to your questions about storing time-series data, number of bands, etc. I focus on standard libraries and best practices common in remote sensing / GIS + machine learning workflows.

Useful libraries and tools / standard workflows

Primary libraries

- **GDAL / OGR**
 - GDAL is the de facto standard for raster processing and format conversion. It supports a very broad set of raster and vector formats and provides command-line tools and Python bindings. waterinag.github.io+1
 - OGR (part of GDAL) handles vector data (e.g. shapefiles), allowing you to read, reproject, rasterize, etc. hydro-informatics.com+1
- **Rasterio (Python)**
 - A user-friendly Pythonic wrapper around GDAL for raster I/O: reading, writing, metadata, easy array access, writing GeoTIFF. automating-gis-processes.github.io+2rasterio.readthedocs.io+2
 - Common in remote sensing + ML workflows because of its integration with NumPy, xarray/rioxarray, etc.
- **GeoPandas / Fiona / Shapely**
 - For vector data manipulation. Fiona for reading/writing shapefiles or other vector formats; GeoPandas builds on Fiona + Shapely for easy vector data analysis and manipulation. [Medium+1](#)

- You can read shapefiles, filter or process geometries/attributes, then rasterize.
- Optionally, high-level toolkits oriented toward ML and geospatial deep learning — e.g. TorchGeo helps integrate geospatial raster datasets with deep learning (PyTorch) pipelines. [arXiv](#)

Example workflows

1. Converting a shapefile (.shp) to a GeoTIFF raster (mask or feature map)

- Use OGR/Fiona (or GeoPandas) to read the shapefile, ensure geometry, CRS.
- Use GDAL (or Rasterio) to rasterize the vector data: set desired pixel size (resolution), define output extent, set target raster dimensions, etc. This produces a TIFF where e.g. a polygon presence/absence or attribute value becomes pixel values. This is standard “vector-to-raster” conversion. [GIS Stack Exchange+2Propeller Aero+2](#)
- In Python with Rasterio, you write the raster data (NumPy array) into a new GeoTIFF, carrying georeferencing, CRS, etc.

2. Converting CSV data (e.g. point observations or gridded values) to GeoTIFF

- If your CSV has coordinates + value (e.g. x, y, value), you can use GDAL’s `gdal.Grid`, or convert to a VRT (virtual raster), or use interpolation to rasterize. In some workflows you produce an intermediate XYZ or VRT before writing GeoTIFF. For example, using GDAL Translate / Rasterize from CSV / XYZ format. [YouTube+1](#)
- Once in raster form (NumPy array), write out via Rasterio or GDAL to GeoTIFF, setting metadata (CRS, transform, nodata, data type).

3. Preprocessing existing GeoTIFF rasters and combining / stacking multiple rasters (e.g., different variables)

- Use Rasterio to open input GeoTIFFs, read arrays, resample / reproject if needed, then write to a single multi-band GeoTIFF. This is common when building multi-variable inputs (e.g. multiple spectral bands, multi-layer features) to feed into a UNet. GDAL supports multi-band GeoTIFF output. [GDAL+2GDAL+2](#)
- You can control compression, tiling/blocking, photometric interpretation, BigTIFF option, etc., via GDAL / Rasterio creation options. [rasterio-spestana.readthedocs.io+1](#)

On bands, time-series, and multi-band TIFFs

- A GeoTIFF raster can have *any number* of bands (within limits) — i.e. it's not restricted to 1 or 3 bands. [GDAL+1](#)
- In principle you can store a time-series as a "stack": one band per time-slice. In a multi-band TIFF, each band is just a 2D raster ($\text{rows} \times \text{cols}$) with consistent georeferencing across bands. That means if your temporal snapshots are co-registered (same extent / resolution / CRS), you could pack them into a single multi-band GeoTIFF, where band 1 = t_0 , band 2 = t_1 , etc. This is often done for multi-date satellite composites or for "data cubes". Support for multi-band is standard. [Carpentries Incubator+2GDAL+2](#)
- The theoretical limit: According to the raster data model of GDAL, the number of bands is limited up to 2,147,483,647 (though in practice a soft limit of 65,536 is more common to avoid excessive RAM/ file size). [GDAL](#)
- Band data types are flexible: you can pick Byte, UInt16, Int16, UInt32, Int32, Float32, Float64, etc. depending on your data requirements. [GDAL+1](#)

Thus — yes — storing time-series as bands in a GeoTIFF is feasible and in fact standard practice in many remote-sensing/data-cube workflows.

Putting it together: example pipeline (shapefile + CSV + raster → input TIFFs for UNet)

Here's a toy pipeline you might adopt (in Python):

```
import geopandas as gpd
import rasterio
from rasterio import features
import pandas as pd
import numpy as np
from osgeo import gdal

# Example: rasterize shapefile into mask
gdf = gpd.read_file('polygons.shp')
```

```
gdf = gdf.to_crs('EPSG:...') # ensure coordinate system
# define raster extent, resolution
transform = rasterio.transform.from_origin(xmin, ymax, pixel_size,
pixel_size)
out_shape = (height, width)

mask = features.rasterize(
    ((geom, 1) for geom in gdf.geometry),
    out_shape=out_shape,
    transform=transform,
    fill=0,
    dtype='uint8'
)

with rasterio.open(
    'mask.tif', 'w',
    driver='GTiff',
    height=out_shape[0], width=out_shape[1],
    count=1, dtype='uint8',
    crs=gdf.crs,
    transform=transform
) as dst:
    dst.write(mask, 1)

# Example: read existing GeoTIFF, stack multiple rasters as bands
with rasterio.open('band1.tif') as src1, rasterio.open('band2.tif') as src2:
    arr1 = src1.read(1)
    arr2 = src2.read(1)
    meta = src1.meta.copy()
    meta.update(count=2)

    with rasterio.open('stacked.tif', 'w', **meta) as dst:
        dst.write(arr1, 1)
        dst.write(arr2, 2)
```

For CSV → GeoTIFF, you could use GDAL via command line (e.g. `gdal.Grid`, `gdal.Translate`), or implement interpolation/resampling in Python (e.g. using `scipy.interpolate`, then rasterize).

Considerations & potential caveats

- Ensure that all datasets (shapefile, CSV coordinate data, rasters) share the *same CRS, resolution, and spatial extent* before stacking or combining into a multi-band TIFF. If they don't, you will need reprojection / resampling.
 - Beware of file size. If you store many time-steps as bands (e.g. hundreds), the resulting TIFF could be huge; consider using compressed GeoTIFF (or even better, cloud-optimized GeoTIFF / tiled formats) or splitting into manageable chunks.
 - When using external tools (e.g. QGIS, SAGA, online converters), double-check that georeference metadata (CRS, transform) is correctly embedded — not all tools preserve spatial metadata reliably. [Geohackweek+1](#)
 - For deep-learning workflows (like UNet), you might also want to consider loading data lazily / using chunks (e.g. with rasterio windows, xarray, or TorchGeo) rather than reading entire giant rasters into memory at once.
-

Summary/Answers to your specific questions

- **Libraries/tools to convert shapefile, CSV, existing GeoTIFFs into TIFFs:** GDAL/OGR, Rasterio, GeoPandas/Fiona/Shapely are standard. These support reading, writing, rasterizing, stacking, reprojection, etc.
 - **Storing time-series data in TIFF:** yes — via multi-band GeoTIFF, using one band per time-slice (assuming consistent spatial grid and georeference).
 - **How many bands can a GeoTIFF have:** theoretically up to ~2.1 billion (but practically soft-limited to 65,536 to avoid resource issues). [GDAL+1](#)
-

If you like, I can **draft a complete reference/well-commented Python script** that takes as input a mix of `.shp`, `.csv`, and existing rasters and outputs a **fully stacked multi-band GeoTIFF** ready for UNet consumption (with optional time-stacking). Do you want me to build that for you now?

You said:

so the primary concerns are having the same crs type, and then rasterizing to ensure each band is properly synced to each pixel through every band?

ChatGPT said:

Yes. For a UNet or any pixel-wise spatial model, **the two non-negotiable requirements** are:

1. All sources must share the same spatial reference (CRS).

Every input layer—shapefiles, CSV-derived rasters, and existing GeoTIFFs—must be in a common CRS before stacking.

If the CRS differs:

- pixel grids won't align,
- extents won't match,
- rasterization will misplaced vector features,
- and the model will learn noise instead of meaningful structure.

Standard operations:

- `gdalwarp` or `rasterio.warp.reproject` for raster reprojection
- `ogr2ogr`, `GeoPandas .to_crs()` for shapefiles and vectors

The goal is a **single CRS**, e.g., EPSG:4326 or a projected CRS like EPSG:3857 or a local UTM zone.

2. Every raster band must be on the same grid: identical transform, extent, and resolution.

This means:

- Same pixel size
- Same raster origin
- Same height × width
- Same [Affine](#) transform
- Same bbox aligned to cell edges

Even if two rasters share a CRS, they can still be offset by fractions of a pixel or use different resolutions. A multi-band GeoTIFF **must** have a perfectly consistent grid across all bands.

If they don't:

- you use GDAL [gdalwarp](#) (resample + align)
 - or Rasterio's [reproject\(\)](#) with a target dataset template.
-

3. Vector data (shp) must be rasterized *after* aligning to the target grid.

You choose a canonical georeferenced grid (typically one of your raster inputs), then rasterize **onto that exact grid**.

For example (Rasterio):

```
mask = features.rasterize(  
    ((geom, value) for geom in gdf.geometry),  
    out_shape=template.shape,  
    transform=template.transform,  
    fill=0,  
    dtype='uint8'
```

)

This ensures:

- each polygon/line/point falls into the exact pixels the model sees,
 - and every derived band perfectly corresponds pixel-by-pixel.
-

4. Multi-band stacking simply becomes concatenation once everything is aligned.

Once every layer:

- is reprojected into the same CRS,
- shares the same resolution + shape,
- uses the same `transform`,

then stacking is trivial:

```
meta = template.meta.copy()
meta.update(count=N)
with rasterio.open("stacked.tif", "w", **meta) as dst:
    for i, arr in enumerate(bands):
        dst.write(arr, i+1)
```

That's the **final input to your UNet**.

Summary

Mandatory conditions for GIS → Deep Learning preprocessing:

1. **Same CRS** for all layers.

2. **Same grid** (pixel size, transform, extent, origin).
3. **Rasterize vectors onto that grid.**
4. **Stack aligned rasters as bands.**

Do these four steps consistently and your UNet will receive clean, spatially correct tensors.

If you want, I can produce a **canonical preprocessing pipeline** that takes arbitrary inputs (.shp, .csv, .tif) and outputs a perfectly aligned multi-band GeoTIFF.

You said:

im wokring with multiband geotiffs. i need to visualize them and validate that they are infact correct. i have heard of qgis, but are there better tools? can qgis read .shp files and tiff files?

ChatGPT said:

QGIS is the standard tool for exactly this task. It reads **GeoTIFF**, **multi-band GeoTIFFs**, **shapefiles (.shp)**, **GeoJSON**, **CSV with geometry**, and essentially every common GIS format. It is built on GDAL, so anything GDAL can read, QGIS can read.

If your goal is to **validate** multi-band rasters (alignment, CRS consistency, pixel values, nodata handling), QGIS is one of the most efficient environments. You can:

- Load each band as its own raster layer or as a stacked composite
- Inspect pixel values with the "Identify Features" tool
- Toggle individual bands on/off
- Visualize RGB composites from arbitrary band combinations
- Check CRS metadata, extents, raster alignment, and nodata regions
- Overlay vector data (shapes, points, polygons) directly on the raster
- Use the raster alignment tool to confirm pixel-perfect registration

Alternatives if you want more specialized workflows:

1. GDAL + Python (Rasterio, rioxarray)

If you need programmatic validation:

- Plot band histograms, pixel distributions
- Check for alignment programmatically (`rasterio.transform`, `rasterio.enums.Resampling`)
- Compute difference maps between bands
- Validate CRS consistency automatically
- Create quick visual previews with matplotlib

This is best for automated QA, but not as convenient for manual inspection.

2. ArcGIS Pro

Commercial, expensive, but very strong for:

- Deep raster analysis
- High-quality visualization
- Time-series rasters
- Advanced symbology

Not necessary unless you already use ESRI software.

3. SNAP (ESA Sentinel Toolbox)

Excellent for satellite datasets, especially Sentinel-1/2 SAR/optical.

Stronger than QGIS for radiometric inspection and band manipulation, but worse for general-purpose vector-raster workflows.

4. Google Earth Engine (GEE)

Good for cloud-based visualization, especially global or large rasters.

But not ideal for local, multi-band TIFF inspection.

5. TerrSet or ENVI

Commercial scientific remote-sensing tools.

Advanced spectral analysis and band math.

Significant cost, overkill unless you do remote-sensing research.

Direct Answer to Your Questions

Can QGIS read .shp files?

Yes. Shapefiles are a native format in QGIS.

Can QGIS read .tif and multi-band GeoTIFFs?

Yes. Fully supported. It can visualize individual bands, composites, and vector overlays.

Recommendation

For pure validation and visualization of multi-band GeoTIFFs + shapefiles, **QGIS is the correct tool**. Nothing is strictly better for your use case unless you need advanced remote-sensing analytics, in which case ENVI or SNAP might be appropriate.

If you also want **automated consistency checks**, use **Python + Rasterio** in combination with QGIS.

You said:

Can qgis show time series data as well? im only seeing static images? how can i under lay a map so that i can confirm the coordinates and the various layers of data that i have? do i need to download a map? how do i import tiffs? i tried to add vector like the .shp files but it fails.

ChatGPT said:

Yes, QGIS can display time-series raster data, but it isn't automatic—you must organize the layers correctly and enable the temporal controller. Below is the minimal workflow to load your TIFFs, load your shapefiles, place them over a basemap, and verify georeferencing.

1. Loading GeoTIFFs Correctly

You **cannot** load **.tif** files using “Add Vector Layer.”

You must use:

Layer → Add Layer → Add Raster Layer

(or click the raster icon in the left toolbar).

Then:

1. Browse to your **.tif**
2. Add it
3. It should appear immediately as a raster layer

If the layer fails to load:

- Make sure it's a **GeoTIFF** (has geotransform + CRS)
- Check that the file is not inside a ZIP

- Verify that GDAL can read it with `gdalinfo yourfile.tif`
-

2. Viewing Multi-Band TIFFs

QGIS automatically detects the number of bands.

Use:

Layer Properties → Symbology → Render type

- *Singleband gray*
- *Multiband color* (choose band 1/2/3 for R/G/B)
- *Band rendering* to cycle each band

This lets you visually inspect band alignment.

3. Adding a Basemap Background

QGIS does *not* ship with basemaps, but you can enable them in one click using the built-in XYZ Tiles.

Add basemap:

1. **Browser Panel → XYZ Tiles → OpenStreetMap (built-in)**
2. Drag it into the map

Or add more basemaps via:

Layer → Data Source Manager → XYZ → New Connection

Use URLs such as:

`https://tile.openstreetmap.org/{z}/{x}/{y}.png`
`https://mt1.google.com/vt/lyrs=s&x={x}&y={y}&z={z}` (Google Satellite)

Now your TIFFs will sit **on top of a real map**, letting you confirm:

- coordinates
 - CRS alignment
 - offsets
 - projections
-

4. Confirming Coordinates and Alignment

For any raster:

Right-click → Properties → Information

- CRS (EPSG:....)
- Pixel size
- Extent (minx, miny, maxx, maxy)
- Geotransform

For shapefiles and rasters with different CRS:

- QGIS will reproject on the fly,
 - But for validation you should use **Raster → Projections → Warp (Reproject)** to create aligned copies.
-

5. Time-Series Raster in QGIS

QGIS supports raster time series, but **you must attach timestamps**.

Workflow:

Option A — If each TIFF represents one date

1. Load all rasters
2. Select them all → Right-click → **Set Temporal Properties**
3. Choose:
 - “**Fixed time range**”
 - Assign start/end times based on filename or metadata
4. Enable **View** → **Panels** → **Temporal Controller**
5. Press **Play** to animate through the series

Option B — If your raster is already a multi-band time cube

QGIS does **not** treat bands as time by default.

You must either:

- split the TIFF into one file per timestamp (`gdal_translate -b n`)
 - or use **Raster** → **Misc** → **Build Virtual Raster (VRT)** to create a band-as-time VRT and map timestamps manually.
-

6. Checking Multiple Layers Are Synchronized

Once loaded:

- Put basemap at the **bottom**
- Put all TIFFs above
- Put shapefiles above those

Use the **Identify Features tool** to click any pixel → see exact coordinates, band values, and vector overlay.

Summary

- Yes, QGIS can display time-series raster if you define temporal metadata.
- Use “**Add Raster Layer**” for TIFFs (not Add Vector).
- Add **XYZ basemaps** to verify geospatial alignment.
- Time-series requires assigning timestamps and enabling the **Temporal Controller**.

If you want, I can give you a minimal step-by-step for your dataset specifically (filenames, band layout, or time format).

Can qgis show time series data as well? im only seeing static images? how can i under lay a map so that i can confirm the coordinates and the various layers of data that i have? do i need to download a map? how do i import tiffs? i tried to add vector like the .shp files but it fails.

ChatGPT said:

Yes, QGIS can display time-series raster data, but it isn't automatic—you must organize the layers correctly and enable the temporal controller. Below is the minimal workflow to load your TIFFs, load your shapefiles, place them over a basemap, and verify georeferencing.

1. Loading GeoTIFFs Correctly

You **cannot** load `.tif` files using “Add Vector Layer.”

You must use:

Layer → Add Layer → Add Raster Layer
(or click the raster icon in the left toolbar).

Then:

1. Browse to your `.tif`
2. Add it
3. It should appear immediately as a raster layer

If the layer fails to load:

- Make sure it's a **GeoTIFF** (has geotransform + CRS)
 - Check that the file is not inside a ZIP
 - Verify that GDAL can read it with `gdalinfo yourfile.tif`
-

2. Viewing Multi-Band TIFFs

QGIS automatically detects the number of bands.

Use:

Layer Properties → Symbology → Render type

- *Singleband gray*
- *Multiband color* (choose band 1/2/3 for R/G/B)
- *Band rendering* to cycle each band

This lets you visually inspect band alignment.

3. Adding a Basemap Background

QGIS does *not* ship with basemaps, but you can enable them in one click using the built-in XYZ Tiles.

Add basemap:

1. **Browser Panel → XYZ Tiles → OpenStreetMap** (built-in)
2. Drag it into the map

Or add more basemaps via:

Layer → Data Source Manager → XYZ → New Connection

Use URLs such as:

`https://tile.openstreetmap.org/{z}/{x}/{y}.png`
`https://mt1.google.com/vt/lyrs=s&x={x}&y={y}&z={z}` (Google Satellite)

Now your TIFFs will sit **on top of a real map**, letting you confirm:

- coordinates
- CRS alignment
- offsets

- projections
-

4. Confirming Coordinates and Alignment

For any raster:

Right-click → Properties → Information

- CRS (EPSG:....)
- Pixel size
- Extent (minx, miny, maxx, maxy)
- Geotransform

For shapefiles and rasters with different CRS:

- QGIS will reproject on the fly,
 - But for validation you should use **Raster → Projections → Warp (Reproject)** to create aligned copies.
-

5. Time-Series Raster in QGIS

QGIS supports raster time series, but **you must attach timestamps**.

Workflow:

Option A — If each TIFF represents one date

1. Load all rasters
2. Select them all → Right-click → **Set Temporal Properties**
3. Choose:

- “**Fixed time range**”
 - Assign start/end times based on filename or metadata
4. Enable **View → Panels → Temporal Controller**
 5. Press **Play** to animate through the series

Option B — If your raster is already a multi-band time cube

QGIS does **not** treat bands as time by default.

You must either:

- split the TIFF into one file per timestamp (`gdal_translate -b n`)
 - or use **Raster → Misc → Build Virtual Raster (VRT)** to create a band-as-time VRT and map timestamps manually.
-

6. Checking Multiple Layers Are Synchronized

Once loaded:

- Put basemap at the **bottom**
- Put all TIFFs above
- Put shapefiles above those

Use the **Identify Features tool** to click any pixel → see exact coordinates, band values, and vector overlay.

Summary

- **Yes**, QGIS can display time-series raster if you define temporal metadata.
- Use “**Add Raster Layer**” for TIFFs (not Add Vector).

- Add **XYZ basemaps** to verify geospatial alignment.
- Time-series requires assigning timestamps and enabling the **Temporal Controller**.

If you want, I can give you a minimal step-by-step for your dataset specifically (filenames, band layout, or time format).

ChatGPT can make mistakes. Check important info.

I will be using a unet on multiband geotiffs and i need to seperate each band. my work flow is to use the last band as the target for the model, that way all environmental data and preceding burn day growth is used as input to determine the last day. it would probably be better to have each day predict the next, but as it is my slow ass computer is take too long to run. the repo code provided this like, but it looks suss. im pretty sure this is actually looking at the first day as the target. which means the model is not working with correct target and inputs. for sanity check this for errors. " target = data[0] #1, H, W inputs = np.concatenate(data[1:]) " " target = data[0] #1, H, W inputs = np.concatenate(data[:-1]) " that is slicing from the beginning to the end so what am i missing?

ChatGPT said:

`data` is almost certainly shaped like:

```
data.shape = (num_bands, H, W)
```

If the bands are ordered chronologically (oldest → newest) then:

- `data[0]` is the first day
- `data[-1]` is the last day

Given that, sanity-checking the two snippets:

1.

```
target = data[0]          # 1, H, W
inputs = np.concatenate(data[1:])
```

This uses **band 0 as the target** and everything *after* band 0 as input.

This is *inverted* relative to your intended workflow (predict last from earlier).

2.

```
target = data[0]          # 1, H, W
inputs = np.concatenate(data[:-1])
```

This still uses **band 0 as the target**, but uses everything *except* the last band as input. It is also **wrong** for the same reason: the target is still the first band.

The difference between the two versions is only which bands appear in the input vector. The target remains incorrect.

What you *actually* want

If “last band” = target, and all preceding bands = input:

```
target = data[-1]          # last band
inputs = np.concatenate(data[:-1])
```

or, more explicitly:

```
target = data[num_bands - 1]
inputs = np.concatenate(data[0:num_bands - 1])
```

Why yours looked “suss”

Because both provided snippets anchor the target at index 0.

Both are slicing the inputs correctly relative to your stated *goal* (either 1: or :-1), but they assume the target is the **first** band. That only makes sense if the repo’s author assumed the first band is the mask.

Sanity check (quick mental model)

Let:

```
bands = [D0, D1, D2, D3]    # D3 = last day
```

Desired:

- inputs: D0, D1, D2

- target: D3

Code should be:

```
inputs = np.concatenate([D0, D1, D2])
target = D3
```

Repo code is doing:

```
inputs = np.concatenate([D1, D2, D3])
target = D0
```

or:

```
inputs = np.concatenate([D0, D1, D2])
target = D0
```

Both are wrong.

Conclusion

You're not missing anything. The repo code is indeed using the **first** band as the target.

Fix it by replacing:

```
target = data[0]
```

with:

```
target = data[-1]
```

If you want, I can rewrite the whole preprocessing block for clarity and safety (including shape assertions and dtype checks).

You said:

where can i find some good implementations of unet architecture? particularly for time series data like ct scans of tumor growth or something. I need to understand how thet inputs work for a different time series project. like what format are the inputs in for these studies and site sources.

ChatGPT said:

If you want to explore good implementations of U-Net (and variants) for medical imaging — especially in contexts close to time-series / longitudinal CT data (e.g. tumor growth over time) — there are several research papers + open-source codebases, and some that even explicitly handle **multiple timepoints** rather than just “single-volume segmentation.” Here’s a breakdown that may help you design your own time-series project.

Key papers / models (including time-series / longitudinal CT / spatio-temporal work)

Work / Model	What they do / Why it matters for longitudinal / CT-scan / time-series use
ImageFlowNet	Forecasts disease progression (e.g. tumor or lesion growth) from longitudinal medical image data. It builds on a UNet-like backbone to learn multiscale features, then uses a neural-ODE/SDE flow framework to model temporal evolution across irregular time-sampling. Applied to datasets including glioblastoma, atrophy, MS, etc. arXiv
LongiUNetDiffWeighting (from MIC-DKFZ)	Extends traditional segmentation approaches by integrating a “Temporal Difference Weighting (TDW)” block: its goal is to explicitly leverage differences between timepoints (not just segment each independently) to improve consistency & capture progression. Meant for longitudinal medical imaging segmentation. GitHub
LA-ResUNet (Longitudinal attention Residual U-Net)	Recently introduced for longitudinal liver tumor segmentation on CT: uses channel & spatial attention plus a “co-segmentation” module to capture cross-temporal correlations and recalibrate features across timepoints. Demonstrates that using timepoint-pair data improves segmentation quality compared to per-timepoint independent segmentation. PubMed

OmniMamba4D	A very recent 2025 work that treats longitudinal CT scans as true 4D data (3 spatial dims + time). Instead of doing separate 3D volumes, it builds spatio-temporal feature blocks (“tetra-orientated Mamba blocks”) to explicitly model spatial and temporal dynamics — useful for analyzing lesion progression/regression, detecting when lesions disappear or appear. arXiv
More “static” but widely used segmentation U-Net/variants (for context)	Typical 3D (or 2D) U-Net and variants still form the backbone for most segmentation tasks — useful when you only have single timepoint scans. E.g. 3D-UNet implementations for lung segmentation. GitHub+2 GitHub+2

Takeaway: While classical U-Net was designed for *single* volume segmentation, recent works are extending it for **longitudinal (time-series) medical imaging** by adding temporal modeling (difference weighting, flow fields, spatio-temporal blocks, attention across timepoints).



Data formats & how input is usually organized

If you're adapting or designing a project, it's important to understand how medical imaging data (especially for CT) is typically stored and fed into models. Common practices from these studies and public datasets:

- Medical imaging commonly uses the **DICOM** standard for raw scans — a series of 2D slices, each as a separate DICOM file. [Wikipedia+2](#)[FiftyOne Documentation+2](#)
- For deep-learning pipelines, especially segmentation or longitudinal analysis, data is often pre-processed and converted to volumetric formats like **NIfTI** ([.nii](#) / [.nii.gz](#)) or other volume formats. This makes it easier to handle as 3D arrays / tensors. [GitHub+2](#)[fdat.uni-tuebingen.de+2](#)
- For longitudinal datasets: each patient will have multiple 3D volumes (baseline, follow-ups). For example, in the dataset described in the “Longitudinal-CT” benchmark (for tumor/metastasis monitoring), each timepoint has a separate NIfTI volume, along with corresponding segmentation masks and lesion metadata (e.g. lesion center-of-gravity). [fdat.uni-tuebingen.de+1](#)
- For models that take time-series/longitudinal inputs (rather than independent volumes), common strategies:

- Concatenate volumes from different timepoints (or their feature maps) as additional channels or separate inputs. Example: in longitudinal segmentation for COVID follow-up CT, a model took as input two 3D volumes (reference + target) to segment the target scan, using the reference segmentation as prior. [arXiv](#)
 - Use temporal-difference modules or flow/ODE frameworks to model changes over time rather than just segment each snapshot. As in LongiUNetDiffWeighting or ImageFlowNet / OmniMamba4D above.
-

Open-source codebases / repos to check (practical starting points)

- **naayem/Medical-CT-Segmentation-DeepLearning** — includes multiple U-Net variants, 3D-UNet, as well as frameworks like nnU-Net and Swin-UNet, oriented to CT segmentation tasks. Good baseline if you just want to try segmentation on volumetric CT data. [GitHub+1](#)
 - **ShawnBIT/UNet-family** — a PyTorch-based collection of original U-Net and many variants (3D U-Net, V-Net, etc.). Good if you want a clean base of “classic U-Nets.” [GitHub](#)
 - **NoeSamaille/unet-lung-segmentation** — specifically demonstrates 3D U-Net segmentation on thoracic CT scans. Shows typical processing of 3D CT volumes. [GitHub](#)
 - **MIC-DKFZ/Longitudinal-Difference-Weighting** — code for temporal-difference weighting block + longitudinal-UNet variants (LongiUNet, LongiUNetDiffWeighting) to handle longitudinal medical image segmentation. [GitHub](#)
 - For cutting-edge 4D longitudinal segmentation: check code / paper of OmniMamba4D (when publicly released; as of now, described in arXiv) — good for modeling spatial + temporal changes together. [arXiv](#)
-

What to watch out for / limitations

- Data availability is often the bottleneck. **Longitudinal, annotated datasets** are rare relative to single-timepoint ones. For instance, the dataset for longitudinal liver tumor

segmentation (used in LA-ResUNet) was newly collected. [PubMed](#)

- Handling time-series introduces complexity: registration between timepoints, normalization across scans (voxel size, intensity), lesion correspondence (tracking same lesion across time), merging/splitting lesions, etc. Many works add explicit modules or pre-processing to handle these.
 - Using standard U-Net on independent timepoints ignores temporal continuity — can cause inconsistency across scans, poor tracking of lesion evolution. For time-series goals (e.g. tumor growth tracking) this is suboptimal.
-

⌚ My recommendation (given your background and project orientation)

Because you're already comfortable coding, and presumably may try something novel (maybe non-medical time-series, but inspired by these methods), here's how I'd proceed:

1. Pull up one of the open-source baseline repos (e.g. [naayem/Medical-CT-Segmentation-DeepLearning](#) or [ShawnBIT/UNet-family](#)) to get familiar with standard 3D U-Net + segmentation pipelines.
 2. Use a small longitudinal dataset (even a 2-timepoint CT dataset like the “Longitudinal-CT” benchmark) to experiment — load volumes as NIfTI, stack timepoints or build time-difference channels, and try naive segmentation to observe issues (mask drift, inconsistencies).
 3. Then experiment with a temporal-aware architecture: try [LongiUNetDiffWeighting](#) or build a custom data pipeline that concatenates timepoints, maybe add attention or temporal-difference mechanisms.
 4. If you want true spatio-temporal modeling, study the design of [OmniMamba4D](#) or [ImageFlowNet](#) — treating scans as 4D (space + time) can be powerful, especially for growth/regression prediction.
-

If you like — I can build a **short list of ~5–10 publicly available longitudinal-CT datasets** (with documentation, file format, # timepoints, etc.) that you could use for prototyping — would make a solid starting resource. Do you want me to build that list for you now?

You said:

so the standard is encoder with convolution then activator. then bottleneck, then decoder with convolution? basically the encoder reduces size and the decoder upsamples after? With each layer im also saving larger features before downsampling to concatiate with the upsampled layers in a symmetrical sort of way? like the sizes must match? i can just use standard nn libraries for this? is there a better alternative than those libraries?

ChatGPT said:

Exactly — that's the core U-Net design. Let's break it down carefully and connect it to what you're asking about:

1. U-Net Architecture Recap

1. Encoder (contracting path):

- Repeated **convolution** → **activation** (e.g., **ReLU**) → **optionally batch norm** → **pooling/downsample**.
- Purpose: extract hierarchical features, reduce spatial resolution while increasing channels.
- Example (PyTorch-style pseudocode):

```
x = conv_block(input_channels, out_channels)
x_down = max_pool2d(x, 2)
```

1.

2. Bottleneck:

- Bottom-most layer where the receptive field is largest.
- Often includes multiple conv layers to encode complex features.

3. Decoder (expanding path):

- Upsample (by transposed conv or interpolation), then **concatenate** with corresponding feature map from encoder.
- Convolution → activation → repeat.

- Purpose: reconstruct spatial resolution, while leveraging high-resolution features saved from encoder (the “skip connections”).

4. Skip Connections:

- Yes, the encoder feature maps are **saved at each level** before downsampling.
- These are concatenated with the decoder’s upsampled feature maps.
- The **spatial dimensions must match** (channels can differ, concatenation along channel dimension).
- If your upsampling slightly differs, cropping or padding is used to match sizes.

5. Output Layer:

- Usually a 1×1 conv to map the final feature map to desired channels (e.g., number of classes).

◆ Why the skip connections matter

- Preserves **high-resolution spatial info** lost during downsampling.
 - Crucial for segmentation tasks where boundaries matter (like tumors in CT scans).
-

2. Implementing in Standard Libraries

Yes, standard deep-learning frameworks are perfectly suited:

- **PyTorch:** most common for research — very flexible for custom U-Net variants.
 - `nn.Conv2d`, `nn.Conv3d` (for volumetric data), `nn.MaxPool2d`, `nn.ConvTranspose2d` or `F.interpolate` for upsampling.
 - There are also community repos implementing ready-to-go U-Net modules for 2D/3D.

- **TensorFlow/Keras:** easy high-level API for prototyping.
 - Layers: `Conv2D`, `Conv3D`, `MaxPooling2D`, `UpSampling2D`, `Concatenate`.
 - Alternatives (less common unless you want specialized performance optimizations):
 - **JAX / Flax / Haiku:** highly flexible, good for custom backprop or research speedups.
 - **MONAI:** PyTorch-based medical imaging library with U-Net, UNETR, nnU-Net, attention variants, etc. (<https://monai.io>).
-

◆ Practical tips

1. Feature map size matching:

- Encoder output before pooling = decoder upsample input size (after upsampling).
- Use padding in conv layers or crop encoder features if needed.

2. Upsampling:

- `ConvTranspose2d` gives learnable upsampling (slightly better) vs. `interpolate` (faster, deterministic).

3. 3D data:

- For volumetric CT or longitudinal scans, just swap `Conv2d` → `Conv3d`, `MaxPool2d` → `MaxPool3d`, etc.
- Time dimension can either be treated as an extra channel (simplest) or as a 4D tensor (space + time) for spatio-temporal architectures.

4. Alternatives / improvements:

- **Attention U-Nets:** add spatial/channel attention in decoder to focus on relevant areas.

- **Residual U-Nets:** add residual connections inside conv blocks for deeper networks.
 - **nnU-Net / MONAI implementations** often outperform hand-rolled U-Nets because they include **normalization, data preprocessing, patch sampling, and learning rate schedules** optimized for medical imaging.
-

✓ Bottom line:

- You can absolutely implement U-Net with standard PyTorch or TensorFlow/Keras.
 - For **medical / volumetric / longitudinal tasks**, MONAI or nnU-Net give a big head start.
 - Skip connections require size matching — usually handled automatically if you keep convolution/pooling consistent.
 - Time-series inputs (like longitudinal CT) can be stacked as channels, or processed with temporal modules if you want temporal dynamics.
-

If you want, I can draw a **visual map of how 3D + timepoint data flows through a U-Net**, including how channels, skips, and upsampling align. That often makes longitudinal data input much clearer.

Do you want me to do that?

You said:

Yes mine is 2d. is there a cool library for visualizing gis data? the model will predict gis change essentially which means the output is also a tiff? that is compared to the target tiff?

ChatGPT said:

Yes — your workflow will be very similar to a standard segmentation/regression task, just in the **GIS / raster domain**. Let's break it down carefully.

1 Input & Output

- **Input:**

- 2D multiband GeoTIFFs (or any raster format) → read into a numpy array of shape `(channels, height, width)` or `(height, width, channels)` depending on the framework.
- Channels could be things like: NDVI, elevation, temperature, or previous burn maps (for wildfire), one channel per “band”/layer.
- **Output:**
 - Same spatial dimensions as input, often a **single-channel 2D raster** representing your prediction (e.g., burn area, flood risk, vegetation growth).
 - Stored as a **GeoTIFF** so it can be overlaid on a map.
 - During training, you **compare the output array to the target array** (ground-truth GeoTIFF) using your loss function. Typical losses: `MSE` (regression), `Dice`, `BCE` (segmentation).

So yes, the predicted raster is **directly comparable to the target raster** pixel-wise.

2 Libraries for working with GIS / raster data

Reading/writing GeoTIFFs

- **rasterio** – Pythonic, standard for raster I/O.

```
import rasterio
with rasterio.open("input.tif") as src:
    img = src.read() # shape: (bands, height, width)
```

-
- **GDAL** – lower-level, very flexible, supports almost all raster formats.
- **xarray + rioxarray** – convenient if you have multiband temporal or multi-dimensional data (like a time series of rasters).

Visualization

- **matplotlib / seaborn** – basic for small arrays.
- **rasterio.plot.show** – quick raster plotting.
- **folium / ipyleaflet** – interactive maps, overlay predictions on basemap.
- **geopandas** – if you have vector overlays (e.g., shapefiles).
- **hvplot + holoviews** – works with xarray for interactive plots.

Example: visualize predicted raster vs ground truth:

```
import rasterio
import matplotlib.pyplot as plt

with rasterio.open("prediction.tif") as pred_src,
rasterio.open("target.tif") as tgt_src:
    pred = pred_src.read(1)
    tgt = tgt_src.read(1)

plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
plt.title("Prediction")
plt.imshow(pred, cmap="viridis")
plt.colorbar()

plt.subplot(1,2,2)
plt.title("Target")
plt.imshow(tgt, cmap="viridis")
plt.colorbar()
plt.show()
```

3 GIS workflow integration with U-Net

1. Read your **multiband input raster** → convert to tensor (PyTorch: (C, H, W)).
2. Pass through **U-Net** → output tensor $(1, H, W)$ (predicted layer).

3. Save prediction tensor as **GeoTIFF** (using rasterio, optionally copying geotransform + CRS from input so it aligns with other GIS data).
4. Compare prediction vs target for **training metrics** (pixel-wise, patch-wise, or other GIS-relevant metrics like IoU, F1, or RMSE).

Optional: if your dataset is large, you can train patch-wise (like classical U-Net) and then **stitch predictions back into a full raster**.

4 Libraries / frameworks that make this easier

- **torchgeo** (<https://torchgeo.readthedocs.io>)
-)
 - High-level framework for segmentation/classification of satellite imagery with U-Net, DeepLab, etc.
- **MONAI** can still be used for raster prediction workflows — just treat rasters like 2D images.

✓ Bottom line:

- Your U-Net predicts a raster of the same shape as input → save as GeoTIFF → compare pixel-wise to target GeoTIFF.
- **rasterio + matplotlib** (or **torchgeo** for more complex datasets) is the standard Python stack for reading, writing, and visualizing GIS data.
- No fundamental change to U-Net; the GIS aspect mainly affects **data I/O, patching, and visualization**.

If you want, I can **draw a diagram showing a 2D multiband GeoTIFF flowing through a U-Net and producing a predicted GeoTIFF**, showing how skip connections, channels, and outputs line up. That's often really clarifying.

Do you want me to do that?

You said:

source good tutorials for qgis, torchgeo, and rastervision

ChatGPT said:

Here are some good tutorial / docs sources for QGIS, TorchGeo, and Raster Vision — including both introductory material and deeper guides.

Recommended Tutorials & Docs

QGIS

- “QGIS Tutorials and Tips” — a widely used collection of step-by-step tutorials for vector & raster operations, map making, data import/export, styling, etc. Great for GIS fundamentals. [QGIS Tutorials](#)
- “GIS Training and Tutorial Resources: QGIS” — a research-library guide summarizing learning resources (beginner → advanced) for QGIS workflows. [UTK Libraries](#)
- A practical guide on combining deep-learning segmentation with QGIS via a plugin (useful if you want to integrate ML outputs back into QGIS): “Deep Learning Image Segmentation in QGIS Using the Deepness Plugin” (Mapscaping article).
[mapscaping.com](#)

TorchGeo

- Official “Introduction to Geospatial Data” tutorial of TorchGeo — covers remote sensing data modalities (multispectral, hyperspectral, raster/vector), data loading, transforms, and how TorchGeo integrates with deep-learning pipelines. [TorchGeo](#)
- “Geospatial deep learning with TorchGeo” — a blog/post describing use-cases, advantages over traditional GIS-only workflows, and how TorchGeo bridges remote sensing imagery with deep learning. [PyTorch](#)
- A hands-on tutorial using TorchGeo + PyTorch Lightning on custom datasets (e.g. satellite / SAR data) — walks through installing TorchGeo, dataset setup, data loading, model training for segmentation. Useful if you want code-level exposure. [Medium](#)
- A practical discussion of data loading challenges in remote sensing & how TorchGeo’s abstractions help — good for understanding pitfalls when working with rasters & geospatial formats. [Sparkgeo](#)

Raster Vision

- Official “Tutorials” page of Raster Vision’s documentation: covers reading raster/vector data, sampling training data, dataset setup, training, prediction, evaluation. Useful to see full pipelines from raw imagery → model → prediction. [Raster Vision+1](#)
- The accompanying Raster Vision documentation (PDF/online) explaining workflow phases (chip generation, training, prediction, evaluation) and use with large imagery sets — helpful for large-scale geospatial ML tasks. [buildmedia.readthedocs.org+1](#)
- A blog post introducing Raster Vision v0.20 — explains updates to the library (multiband support, integration with deep-learning backends, using Raster Vision as a library instead of black-box framework). Useful to understand design philosophy and practical usage. [Medium](#)

⌚ Which to Use for What

Goal / Task	Suggested Tool / Tutorial
Learn GIS basics (raster/vector styling, map visualizing, coordinate reference systems, overlays)	QGIS tutorials (QGIS Tutorials and Tips; QGIS Training Resource)
Combine ML segmentation/regression with GIS — visualize results, integrate with maps/shapefiles	QGIS + deep-learning plugin approach (Deepness guide)
Build ML-ready geospatial pipelines using Python + PyTorch	TorchGeo tutorials + blog posts
Work with large-scale satellite / aerial imagery, multiband rasters, segmentation / detection tasks on geospatial imagery	Raster Vision official tutorials + docs + blog

Because your interest is in **predicting GIS changes** and likely working with multi-channel raster data (e.g. GeoTIFFs), **TorchGeo** and **Raster Vision** are likely most relevant to you. Then you can **use QGIS for visualization, validation, and mapping** of results.

If you like, I can also include **3–5 real example GitHub repos** (with code) that use these tools for GIS + deep learning — that way you can see end-to-end pipelines (data load → model → prediction → output geo-raster).

) — excellent for longitudinal / volumetric medical data.

nnU-Net framework: automates patching, normalization, data augmentation, and trains SOTA segmentation networks with minimal effort (<https://github.com/MIC-DKFZ/nnUNet>

)

- Built for remote sensing / GIS data in PyTorch. Handles raster reading, patching, and geospatial-aware datasets.

rastervision (<https://docs.rastervision.io>

popular loss functions for unet. unet and time series. give sources.

ChatGPT said:

For **UNet**, especially in segmentation tasks (including spatial-temporal tasks like wildfire spread prediction or medical imaging), the choice of loss function is critical. Here's a breakdown:

1. Popular Loss Functions for UNet

a. Dice Loss

- **Definition:** Measures overlap between predicted mask and ground truth.
- **Formula:**

$$\text{Dice Loss} = 1 - \frac{2 \sum_i p_i g_i}{\sum_i p_i + \sum_i g_i}$$

where p_i is the predicted probability, g_i is the ground truth.

- **Pros:** Handles class imbalance well, focuses on shape and overlap.
 - **Cons:** Can be unstable when both prediction and target are empty.
 - **Use-case:** Segmentation with sparse positive regions (e.g., small wildfires, tumors).
-

b. Cross-Entropy Loss (Binary / Categorical)

- **Definition:** Standard pixel-wise classification loss.
 - **Pros:** Stable, well-understood.
 - **Cons:** Can underperform with severe class imbalance.
 - **Use-case:** When foreground and background are roughly balanced or when combined with Dice (hybrid loss).
-

c. Focal Loss

- **Definition:** Modifies cross-entropy to focus on hard-to-classify pixels.

$$\text{FL}(pt) = -\alpha(1-pt)\log(pt) \quad \text{FL}(p_t) = -\alpha_t(1-p_t)^\gamma$$

- **Pros:** Good for highly imbalanced datasets.
 - **Cons:** Hyperparameter tuning (α, γ) needed.
 - **Use-case:** Sparse wildfire pixels in a large raster.
-

d. Tversky Loss

- **Definition:** Generalization of Dice Loss with tunable false positive/negative weighting.

$$\text{Tversky Index} = \frac{\sum_i p_i g_i + \alpha \sum_i p_i (1-g_i) + \beta \sum_i (1-p_i) g_i}{\sum_i p_i g_i + \alpha \sum_i p_i (1-g_i) + \beta \sum_i (1-p_i) g_i}$$

- **Pros:** Flexible for controlling FP/FN trade-off.
 - **Cons:** More hyperparameters to tune.
 - **Use-case:** Wildfire growth prediction where overprediction is worse than underprediction (or vice versa).
-

e. Mean Squared Error (MSE) / L1 Loss

- **Use-case:** When UNet is used for **regression** tasks (e.g., predicting continuous environmental variables per pixel or next-day burned area in time series).
-

f. Hybrid Losses

- Combining **Dice + Cross-Entropy** or **Tversky + Focal** often improves convergence and balances class imbalance with pixel-wise accuracy.
-

2. UNet for Time Series / Spatio-Temporal Data

- **Problem:** Standard UNet is 2D (spatial). Time series adds temporal dimension.
 - **Approaches:**
 1. **3D UNet:** Treat time as a third dimension ($C \times H \times W \times TC$ \times $H \times W \times TC$).
 2. **UNet + ConvLSTM:** Encode spatial features per timestep, then use ConvLSTM to model temporal evolution.
 3. **Sliding window approach:** Feed consecutive time slices as multi-channel input to 2D UNet (like you mentioned using past days as input channels).
 - **Loss functions:** Same as above. If predicting **continuous growth**, MSE or MAE might replace Dice. If predicting **binary masks** (burned/not burned), Dice or Tversky works well.
-

3. References / Sources

1. Original UNet:

Ronneberger et al., *U-Net: Convolutional Networks for Biomedical Image Segmentation*, MICCAI 2015. <https://arxiv.org/abs/1505.04597>

Dice / Tversky Loss:

Salehi et al., *Tversky Loss for Image Segmentation using 3D Fully Convolutional Networks*, 2017. <https://arxiv.org/abs/1706.05721>

Focal Loss:

Lin et al., *Focal Loss for Dense Object Detection*, 2017. <https://arxiv.org/abs/1708.02002>

UNet + Time Series / Spatio-Temporal:

- Shi et al., *Convolutional LSTM Network: A Machine Learning Approach for Precipitation Nowcasting*, NIPS 2015. <https://arxiv.org/abs/1506.04214>

Zhang et al., *Spatio-Temporal UNet for Traffic Flow Prediction*, 2021.
<https://arxiv.org/abs/2004.00573>

If you want, I can make a **table summarizing which loss functions are best for wildfire/earth observation datasets with UNet** including pros, cons, and expected F1 behavior. It'll make choosing the right one much faster. Do you want me to do that?